

SUPUN ABEYSINGHE

🏠 supun.online ✉ supunabeysinghe@gmail.com 🌐 github.com/supunab 📄 in/supunabeysinghe
📍 West Lafayette, IN, USA 📞 +1 (765)-775-8679

EDUCATION

Purdue University

Ph.D. in Computer Science - GPA 4.0/4.0

Advisor: Prof. Tiark Rompf

Relevant Courses: Operating Systems, Algorithms, Compilers, Database Systems, Distributed Systems

West Lafayette, IN

Expected Dec '23

University of Moratuwa

B.Sc Engineering (Hons) (Computer Science & Engineering)

First Class Honours - GPA 4.06/4.20

Sri Lanka

Dec '18

EXPERIENCE

Research Intern, Microsoft

Sep 2022 - Dec 2022

- Developed a compiler prototype to generate specialized ONNXRuntime (ORT) custom ops from ONNX models, utilizing efficient specialized kernels (CUTLASS). Demonstrated significant performance improvements for critical models such as GPT2 and BERT compared to the standard ORT implementation.

Software Engineer Intern, SambaNova Systems

May 2022 - July 2022

- Spearheaded optimization efforts on a critical analysis pass in the compiler stack (MLIR) to reduce resource consumption. Achieved a significant increase in resource utilization and successfully integrated the optimizations into production

Graduate Research Assistant, Purdue University

May 2020 - Present

- Led three research projects focused on applying runtime code generation techniques to develop high-performance data analytics (SQL and Datalog) and ML systems. Successfully submitted three papers for publication, with one already accepted at SIGMOD '22 and others currently under review.

Research Engineer, WSO2 Inc. Sri Lanka

Jan 2019 - July 2019

- Contributed to a research project focused on leveraging ML techniques to dynamically auto-tune server configurations. Co-authored a paper published in ISCC '19, which presented novel findings on the effectiveness of the approach

Research Intern, StatNLP Lab, SUTD Singapore

Jun 2017 - Dec 2017

- Was responsible for testing and fixing bugs of the StatNLP framework by implementing several traditional ML models

SELECTED PROJECTS

- Optimizing End-to-end Data Science Pipelines:** Leveraging generative programming techniques to accelerate combined data processing and ML workloads by constructing common intermediate (IR) layer integrations. Observed **speedups up to 20x** in end-to-end performance (Scala, C++, CUDA)
- Efficient Incrementalization of SQL Queries with Nested Aggregates:** Building novel tree-based index structures to improve the incrementalization efficiency of nested-aggregate queries by **up to 1000x** over the state-of-the-art
- Building Efficient and Expressive Datalog Systems:** Building a novel Datalog compiler for declarative program analysis, leveraging generative programming techniques to achieve up to **10x speedups** compared to state-of-the-art engines such as Souffle. Introduced a novel query execution mechanism based on worst-case optimal joins, along with associated index structures, which significantly improve performance and scalability
- Server Parameter Auto-tuning using Machine Learning (2019):** Leveraging ML based optimization techniques to dynamically auto-tune server parameters to enhance runtime performance (Java, Python, PyTorch)
- Social Media Analytics Platform (2018):** A platform for automatically extracting information related to a particular entity (e.g., restaurant reviews from multiple sources) and performing various types of analysis (emotion detection, aspect-based sentiment analysis, trending topics and evolution of them, etc.) (Python, PyTorch)

SELECTED COURSE PROJECTS

1. **Graph Query Compilation:** Extending relational query compilation techniques based on generative programming to support compilation of graph queries, achieving **an order of magnitude speedup** compared to interpreted engines (Scala, C++)
2. **Sharded, Distributed Key-Value Store:** Implemented a sharded, distributed KV store using Paxos algorithm for replication and support for transactions using Two-phase Commit (**Java, Distributed Systems**)
3. **University C Compiler:** A compiler for a C-like programming language that uses LLVM as an IR. Written multiple LLVM passes for compiler optimizations (**C++, LLVM**)

TECHNICAL SKILLS

- **Programming Languages:** Scala, Java, Python, C, C++
- **Systems and Libraries:** Spark, Flink, PyTorch, Tensorflow, CUDA, LLVM, MLIR;
Prior Experience - Spring Boot, Angular, Node.js

SELECTED OPEN SOURCE CONTRIBUTIONS

Lantern - Deep Learning Framework

- Implemented custom CUDA kernels for several Deep Learning operators (Softmax, Embedding layer, etc.), added support for modern deep learning models like Transformers, etc.

Lightweight Modular Staging (LMS) - A Compiler Framework

- Several key feature additions (e.g., lambda lifting support for top level functions for the C codegen backend) and general maintenance

SELECTED PUBLICATIONS

1. **Efficient Incrementalization of Correlated Nested Aggregate Queries using Relative Partial Aggregate Indexes (RPAI)**
Supun Abeysinghe, Qiyang He, Tiark Rompf
(To appear at SIGMOD '22)
2. **Architecting Intermediate Layers for Efficient Composition in End-to-End Data Science Pipelines**
Supun Abeysinghe, Fei Wang, Gregory Essertel, Tiark Rompf
(in submission)
3. **Efficient and Expressive Datalog Engines: A Generative Programming Approach**
Supun Abeysinghe, Tiark Rompf
(in submission)
4. **ADAPT-T: An Adaptive Algorithm for Auto-Tuning Worker Thread Pool Size in Application Servers**
Nilushan Costa, Malith Jayasinghe, Ajantha Atukorale, Supun Abeysinghe, Srinath Perera, Isuru Perera
At International Symposium on Computers and Communications (ISCC), Barcelona, Spain (2019)
5. **Sentylic at IEST 2018: Gated Recurrent Neural Network and Capsule Network Based Approach for Implicit Emotion Detection**
Prabod Rathnayaka, Supun Abeysinghe, Chamod Samarajeewa, Isura Manchanayake, Malaka Walpola
At Workshop on Computational Approaches to Subjectivity and Sentiment Analysis (held with EMNLP) (2018)

SELECTED AWARDS

- **Competitive Programming:** IEEEExtreme v12.0 (2018) - Country Rank 1st Global Rank 60th (top 2%), IEEEExtreme v10.0 (2016) - Country Rank 3rd Global Rank 106th (top 5%)
- 5th place - Implicit Emotion Detection Shared Task, Organized as part of WASSA 2018 at EMNLP 2018
- **Dean's List Award:** Dean's list award for all semesters (should achieve a semester GPA of 3.8 or above to be included in the Dean's List) (in undergraduate)